

YASH DESHMUKH

+91-7247662800 ◇ Bengaluru, IN ◇ yashdeshmukh7@gmail.com ◇ [linkedin](#) ◇ [github](#)

SUMMARY

Senior Backend Engineer with 5 years of experience building distributed systems that process 1M+ GPS telemetry events/min, 680GB/day, and sustain 99.99% availability. Specializes in Java, Spring Boot, Kafka, Cassandra, Redis, and AWS across edtech, logistics, and supply chain. Proven record of owning end-to-end architecture, mentoring engineers, and driving design reviews from ambiguous requirements to production-grade systems.

CORE COMPETENCIES

Cloud & Architecture	AWS, GCP, S3, Redshift, Data Lakes, Microservices, Distributed Systems, High Availability, Scalability, Reliability Engineering, Observability
Backend Engineering	Java, Python, Spring Boot, Spring Batch, Multithreading, Distributed Locking, Event-Driven Architecture, REST APIs, SDLC
Databases & Storage	MySQL, PostgreSQL, Cassandra, Redis, Amazon Redshift, SQLite, Athena, Data Warehousing, CDC Pipelines
Messaging & Automation	Apache Kafka, AWS SQS, Apache Airflow, ActiveMQ, Kafka Connect, Background Jobs, Batch Processing

EXPERIENCE

Software Developer 2 - Superset — Java, Spring Boot, MySQL, Cassandra, AWS, SQS — Nov 2024 - Present

- Architected an asynchronous SQL assessment engine supporting 200MB SQLite datasets with connection-pool resource limits and decoupled Judge0 worker evaluation; scaled to 1,000 concurrent assessments across 8 universities and 2 enterprise customers, enabling high-margin data-engineering hiring tracks.
- Reduced operations team proctoring review time by 50% by building a violation-review framework supporting bulk verification of 100 attempts per request, configurable result-access gating, and band classification for high-stakes campus hiring.
- Maintained contracted SLA commitments for client-facing dashboards by engineering Airflow-driven Python operator pipelines across 17 Redshift tables and a CDC DAG that refreshed high-traffic aggregate tables every 5 minutes, replacing a daily batch.
- Eliminated 100% of lock-wait timeout failures during peak assessment scoring by refactoring MySQL transaction boundaries and decoupling read-intent locks from update operations; mentored 3 engineers and introduced a structured on-call workflow.

Software Developer 2 - Blackbuck — Java, Redis, Kafka, New Relic, MySQL — May 2023 - Oct 2024

- Owned end-to-end architecture of a real-time Kafka telemetry ingestion pipeline processing 1M+ GPS events; resolved consumer-group rebalancing failures causing data loss during peak bursts and delivered a fuel consumption dashboard with technician SOP workflow.
- Reduced critical API latency by 87% from 300ms to 40ms by profiling slow transactions with New Relic and refactoring database access paths in the shipment tracking stack.
- Delivered full-stack AutoPay workflows using Razorpay eNACH for automated mandate flows and restructured database access patterns and renewal-processing logic for reliability; contributed to a 40% improvement in GPS subscription renewal rates.
- Mentored 5 engineers, introduced a formal design-review document process, and identified Redis and MySQL throughput bottlenecks through MBean-based profiling with targeted optimizations.

Software Developer 1 - Blackbuck — Java, Spring Boot, AWS, Data Lakes, Python — Jun 2021 - May 2023

- Maintained 99.99% availability on a fulfillment platform processing 1M+ RPM by leading root-cause analysis and systematic resolution of 200+ production incidents.

- Migrated high-throughput location-tracking from MongoDB to Cassandra, redesigning wide-column partition and clustering keys to absorb write-heavy telemetry workloads and scale warehouse aggregation pipelines.
- Delivered a 680GB/day cloud data platform by engineering Cassandra-to-AWS data lake pipelines using Kafka Connect and Amazon Athena; executed zero-downtime migration of GPS and telematics services from AWS to GCP.

EDUCATION

Bachelor of Technology - NIT Warangal
Electronics And Communications Engineering

2017 - 2021
CGPA - 7.89/10